

```

# Data manipulation and analysis
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')

# Statistical tests and time series analysis
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# Machine learning
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Set plotting style
plt.style.use('seaborn-v0_8')
plt.rcParams['figure.figsize'] = (12, 8)

# Load the NIFTY-50 dataset
df = pd.read_csv('nifty50.csv')

# Display basic information about the dataset
print("Dataset Shape:", df.shape)
print("\nColumn Names:")
print(df.columns.tolist())
print("\nFirst 5 rows:")
print(df.head())
print("\nDataset Info:")
print(df.info())
print("\nBasic Statistics:")
print(df.describe())

Dataset Shape: (249, 7)

Column Names:
['Date', 'Open', 'High', 'Low', 'Close', 'Shares Traded', 'Turnover (?
Cr)']

First 5 rows:

```

Date	Open	High	Low	Close	Shares Traded	Turnover (? Cr)
2000-01-03	100.00	100.00	100.00	100.00	1000000	100000000.00
2000-01-04	100.00	100.00	100.00	100.00	1000000	100000000.00
2000-01-05	100.00	100.00	100.00	100.00	1000000	100000000.00
2000-01-06	100.00	100.00	100.00	100.00	1000000	100000000.00
2000-01-09	100.00	100.00	100.00	100.00	1000000	100000000.00

0	27-Jun-24	23881.55	24087.45	23805.4	24044.50	515227010
1	28-Jun-24	24085.90	24174.00	23985.8	24010.60	354779832
2	01-Jul-24	23992.95	24164.00	23992.7	24141.95	242468081
3	02-Jul-24	24228.75	24236.35	24056.4	24123.85	309629240
4	03-Jul-24	24291.75	24309.15	24207.1	24286.50	289201551

	Turnover (? Cr)
0	61216.71
1	41242.87
2	28204.06
3	34838.65
4	36661.17

#### Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 249 entries, 0 to 248
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Date	249 non-null	object
1	Open	245 non-null	float64
2	High	247 non-null	float64
3	Low	248 non-null	float64
4	Close	247 non-null	float64
5	Shares Traded	249 non-null	int64
6	Turnover (? Cr)	247 non-null	float64

```
dtypes: float64(5), int64(1), object(1)
```

```
memory usage: 13.7+ KB
```

```
None
```

#### Basic Statistics:

	Open	High	Low	Close	Shares
Traded \					
count	245.000000	247.000000	248.000000	247.000000	
mean	24171.996122	24294.145344	24050.545766	24172.141296	
std	875.759548	868.484538	877.796724	872.667163	
min	21758.400000	22105.050000	21743.650000	22082.650000	
25%	23542.150000	23689.675000	23430.912500	23529.600000	
50%	24328.850000	24402.650000	24182.850000	24323.850000	
75%	24812.600000	24946.875000	24688.725000	24818.300000	
max	26248.250000	26277.350000	26151.400000	26216.050000	

	Turnover (? Cr)
count	247.000000
mean	30385.644737
std	8809.522852
min	3348.450000
25%	25199.570000
50%	28871.350000
75%	33792.955000
max	85975.810000

## Data Exploration

### 1.1 Missing Values Analysis and Imputation Techniques

```
print("Missing Values Analysis:")
print("=" * 50)
missing_values = df.isnull().sum()
missing_percentage = (missing_values / len(df)) * 100

missing_df = pd.DataFrame({
    'Column': missing_values.index,
    'Missing Count': missing_values.values,
    'Missing Percentage': missing_percentage.values
})

print(missing_df[missing_df['Missing Count'] > 0])

print(f"\nTotal missing values: {df.isnull().sum().sum()}")
print(f"Percentage of complete cases:
{(df.dropna().shape[0]/df.shape[0])*100:.2f}%")
```

Missing Values Analysis:

```
=====
      Column  Missing Count  Missing Percentage
1         Open             4           1.606426
2         High             2           0.803213
3         Low              1           0.401606
4         Close            2           0.803213
6  Turnover (? Cr)         2           0.803213
```

Total missing values: 11

Percentage of complete cases: 95.58%

### Imputation Techniques for NIFTY-50 Time Series Data

**1. Forward Fill:** Uses the last valid observation to fill missing values

- **Suitable for:** Stock prices where values change gradually
- **Assumption:** Recent price is the best estimate for missing value

**2. Backward Fill:** Uses the next valid observation to fill missing values

- **Suitable for:** When future information is available and relevant

**3. Linear Interpolation:** Estimates missing values using linear relationship between known points

- **Suitable for:** Time series with smooth trends
- **Best for:** Stock price data with consistent trends

**4. Moving Average:** Uses average of surrounding values

- **Suitable for:** Smoothing out short-term fluctuations
- **Parameters:** Window size determines smoothing level

**5. Seasonal Interpolation:** Considers seasonal patterns in the data

- **Suitable for:** Data with clear seasonal patterns
- **Best for:** Long-term stock data with recurring patterns

```
# Create a copy for demonstration
df_demo = df.copy()

#selected Linear Interpolation
df_interpolate = df_demo.interpolate(method='linear')

df_clean = df_interpolate.copy()
print(f"Selected method: Linear Interpolation")
print(f"Remaining missing values: {df_clean.isnull().sum().sum()}")

Selected method: Linear Interpolation
Remaining missing values: 0
```

## 1.2 Data Preprocessing and Visualization

```
# Assuming the dataset has a date column, convert it to datetime
date_columns = [col for col in df_clean.columns if 'date' in
col.lower() or 'time' in col.lower()]
if date_columns:
    df_clean[date_columns[0]] =
pd.to_datetime(df_clean[date_columns[0]])
    df_clean.set_index(date_columns[0], inplace=True)
    print(f"Date column '{date_columns[0]}' converted to datetime and
set as index")

# Display data types and sample data
print("\nData Types:")
```

```
print(df_clean.dtypes)
print("\nSample Data:")
print(df_clean.head())
```

Date column 'Date' converted to datetime and set as index

Data Types:

```
Open          float64
High          float64
Low           float64
Close         float64
Shares Traded int64
Turnover (? Cr) float64
dtype: object
```

Sample Data:

	Open	High	Low	Close	Shares Traded \
Date					
2024-06-27	23881.55	24087.45	23805.4	24044.50	515227010
2024-06-28	24085.90	24174.00	23985.8	24010.60	354779832
2024-07-01	23992.95	24164.00	23992.7	24141.95	242468081
2024-07-02	24228.75	24236.35	24056.4	24123.85	309629240
2024-07-03	24291.75	24309.15	24207.1	24286.50	289201551

Turnover (? Cr)

Date	
2024-06-27	61216.71
2024-06-28	41242.87
2024-07-01	28204.06
2024-07-02	34838.65
2024-07-03	36661.17

*# Get numeric columns for visualization*

```
numeric_cols =
df_clean.select_dtypes(include=[np.number]).columns.tolist()
```

*# 1. Line Charts - Time Series Analysis*

```
fig, axes = plt.subplots(2, 2, figsize=(20, 12))
fig.suptitle('NIFTY-50 Time Series Analysis - Line Charts',
fontsize=16, y=1.02)
```

*# Plot first 4 numeric columns as line charts*

```
for i, col in enumerate(numeric_cols[:4]):
    row, col_idx = i // 2, i % 2
    axes[row, col_idx].plot(df_clean.index, df_clean[col],
linewidth=1.5, alpha=0.8)
    axes[row, col_idx].set_title(f'{col} Over Time', fontsize=12,
fontweight='bold')
    axes[row, col_idx].set_xlabel('Date')
```

```

axes[row, col_idx].set_ylabel(col)
axes[row, col_idx].grid(True, alpha=0.3)
axes[row, col_idx].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

# 2. Scatter Plots - Correlation Analysis
if len(numeric_cols) >= 2:
    fig, axes = plt.subplots(2, 2, figsize=(15, 12))
    fig.suptitle('NIFTY-50 Scatter Plot Analysis', fontsize=16,
y=1.02)

    combinations = [(0,1), (0,2), (1,2), (1,3)] if len(numeric_cols)
>= 4 else [(0,1)] * 4

    for i, (x_idx, y_idx) in enumerate(combinations):
        if x_idx < len(numeric_cols) and y_idx < len(numeric_cols):
            row, col_idx = i // 2, i % 2
            x_col, y_col = numeric_cols[x_idx], numeric_cols[y_idx]

            axes[row, col_idx].scatter(df_clean[x_col],
df_clean[y_col], alpha=0.6, s=30)
            axes[row, col_idx].set_xlabel(x_col)
            axes[row, col_idx].set_ylabel(y_col)
            axes[row, col_idx].set_title(f'{x_col} vs {y_col}')
            axes[row, col_idx].grid(True, alpha=0.3)

            # Add correlation coefficient
            corr = df_clean[x_col].corr(df_clean[y_col])
            axes[row, col_idx].text(0.05, 0.95, f'Corr: {corr:.3f}',
transform=axes[row,
col_idx].transAxes,
bbox=dict(boxstyle='round',
facecolor='white', alpha=0.8))

        plt.tight_layout()
        plt.show()

# 3. Histograms - Distribution Analysis
n_cols = min(4, len(numeric_cols))
fig, axes = plt.subplots(2, 2, figsize=(15, 10))
fig.suptitle('NIFTY-50 Distribution Analysis - Histograms',
fontsize=16, y=1.02)

for i in range(4):
    row, col_idx = i // 2, i % 2
    if i < len(numeric_cols):
        col = numeric_cols[i]
        axes[row, col_idx].hist(df_clean[col], bins=50, alpha=0.7,

```

```

edgecolor='black', linewidth=0.5)
    axes[row, col_idx].set_title(f'Distribution of {col}',
fontweight='bold')
    axes[row, col_idx].set_xlabel(col)
    axes[row, col_idx].set_ylabel('Frequency')
    axes[row, col_idx].grid(True, alpha=0.3)

    # Add statistics
    mean_val = df_clean[col].mean()
    std_val = df_clean[col].std()
    axes[row, col_idx].axvline(mean_val, color='red',
linestyle='--', alpha=0.8, label=f'Mean: {mean_val:.2f}')
    axes[row, col_idx].legend()
else:
    axes[row, col_idx].axis('off')


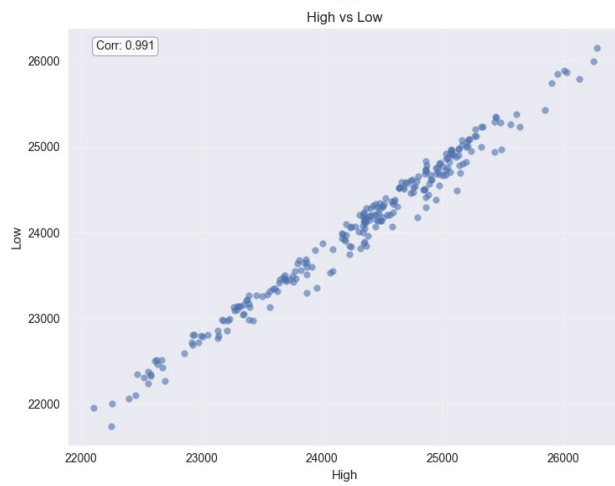

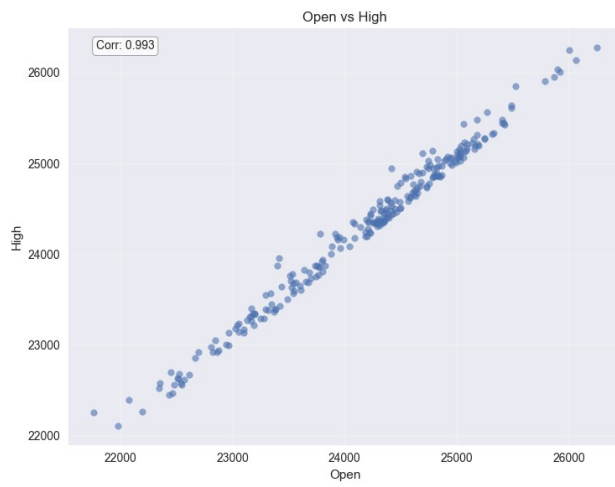
plt.tight_layout()
plt.show()

```

NIFTY-50 Time Series Analysis - Line Charts

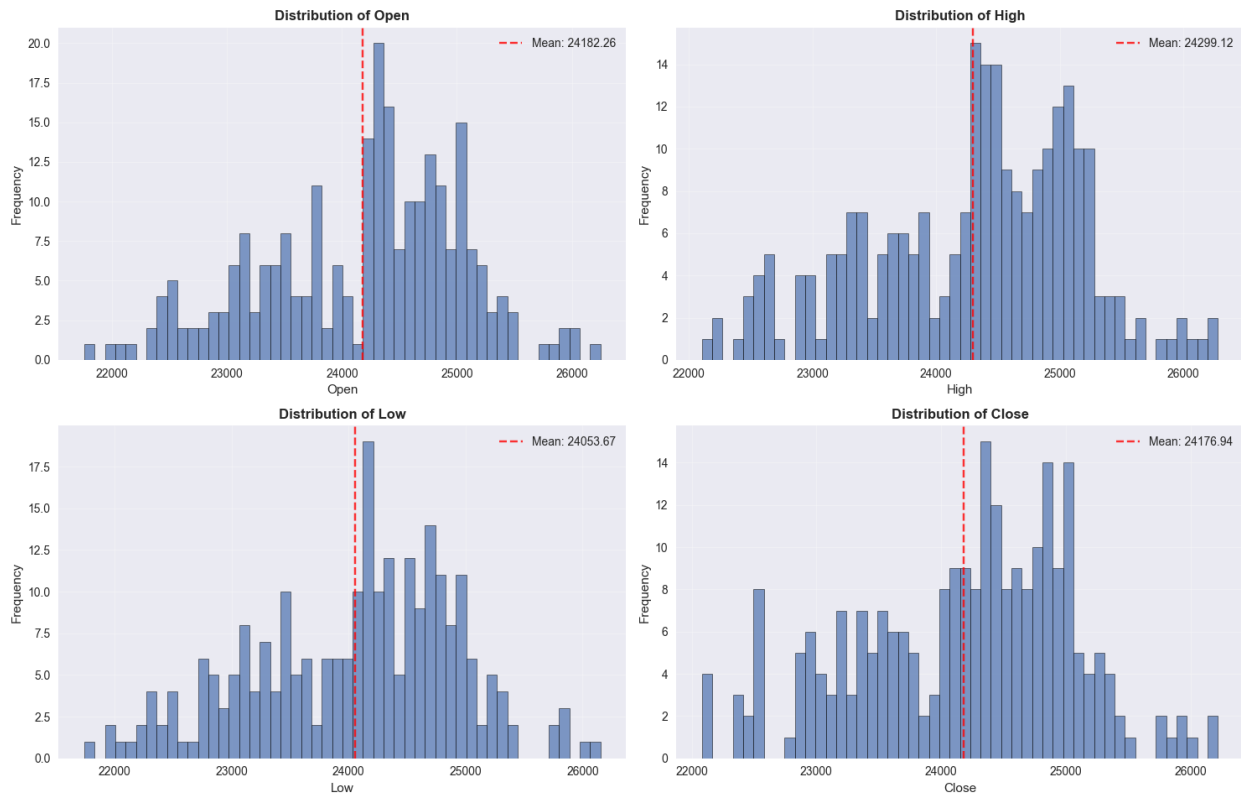


# NIFTY-50 Scatter Plot Analysis





## NIFTY-50 Distribution Analysis - Histograms



## 2.1 Turnover Prediction Model

**Objective:** Build the best model to predict "Turnover" by identifying suitable variables and providing justification.

```
# Correlation Analysis for Feature Selection
print("\n" + "="*50)
print("CORRELATION ANALYSIS")
print("="*50)

# Define modeling dataframe and features
df_model = df_clean.copy()
target_col = "Turnover (? Cr)"
all_features = [col for col in
df_model.select_dtypes(include=[np.number]).columns if col !=
target_col]

# Calculate correlations with target variable
correlations =
df_model[all_features].corrwith(df_model[target_col]).abs().sort_value
s(ascending=False)

print("Top 15 features correlated with Turnover:")
```

```

print(correlations.head(15))

# Select top features based on correlation
top_features = correlations.head(15).index.tolist()

# Correlation heatmap for top features
plt.figure(figsize=(12, 10))
correlation_matrix = df_model[top_features + [target_col]].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, fmt='.2f', cbar_kws={'label': 'Correlation
Coefficient'})
plt.title('Correlation Matrix - Top Features vs Turnover')
plt.tight_layout()
plt.show()

print(f"\nSelected features for modeling: {top_features}")

```

```

=====
CORRELATION ANALYSIS
=====


```

```

Top 15 features correlated with Turnover:

```

Shares Traded	0.812498
High	0.189621
Close	0.171003
Open	0.162373
Low	0.138871
dtype: float64	



Selected features for modeling: ['Shares Traded', 'High', 'Close', 'Open', 'Low']

*# Model Training and Evaluation*

```
print("\n" + "="*50)
```

```
print("MODEL TRAINING AND EVALUATION")
```

```
print("="*50)
```

*# Prepare data for modeling*

```
X = df_model[top_features]
```

```
y = df_model[target_col]
```

*# Split data (time series split - use earlier data for training)*

```
split_idx = int(len(df_model) * 0.8)
```

```
X_train, X_test = X[:split_idx], X[split_idx:]
```

```
y_train, y_test = y[:split_idx], y[split_idx:]
```

```

print(f"Training set size: {X_train.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")

# Initialize models
models = {
    'Linear Regression': LinearRegression(),
    'Ridge Regression': Ridge(alpha=1.0),
    'Lasso Regression': Lasso(alpha=1.0),
    'Random Forest': RandomForestRegressor(n_estimators=100,
random_state=42)
}

# Train and evaluate models
results = {}

for name, model in models.items():
    # Train model
    model.fit(X_train, y_train)

    # Make predictions
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    # Calculate metrics
    train_mae = mean_absolute_error(y_train, y_pred_train)
    test_mae = mean_absolute_error(y_test, y_pred_test)
    train_rmse = np.sqrt(mean_squared_error(y_train, y_pred_train))
    test_rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
    train_r2 = r2_score(y_train, y_pred_train)
    test_r2 = r2_score(y_test, y_pred_test)

    results[name] = {
        'Train MAE': train_mae,
        'Test MAE': test_mae,
        'Train RMSE': train_rmse,
        'Test RMSE': test_rmse,
        'Train R²': train_r2,
        'Test R²': test_r2,
        'Model': model,
        'Predictions': y_pred_test
    }

# Display results
results_df = pd.DataFrame({name: {metric: values[metric] for metric in
['Train MAE', 'Test MAE', 'Train RMSE', 'Test RMSE', 'Train R²', 'Test
R²']}}
                        for name, values in results.items()}).T

print("\nModel Performance Comparison:")

```

```

print(results_df.round(4))

# Find best model based on Test R2
best_model_name = results_df['Test R2'].idxmax()
best_model = results[best_model_name]['Model']

print(f"\nBest Model: {best_model_name}")
print(f"Test R2 Score: {results_df.loc[best_model_name, 'Test R2']:.4f}")
print(f"Test RMSE: {results_df.loc[best_model_name, 'Test RMSE']:.4f}")

=====
MODEL TRAINING AND EVALUATION
=====
Training set size: 199
Test set size: 50

Model Performance Comparison:

```

	Train MAE	Test MAE	Train RMSE	Test RMSE	Train R <sup>2</sup>
Linear Regression	2385.2157	7893.3074	3881.2347	9511.4805	0.8186
Ridge Regression	2385.2157	7893.3073	3881.2347	9511.4805	0.8186
Lasso Regression	2385.8054	7894.9474	3881.2440	9511.7305	0.8186
Random Forest	1016.5678	7571.0970	1875.7307	9454.4696	0.9576

```


```

	Test R <sup>2</sup>
Linear Regression	-0.5446
Ridge Regression	-0.5446
Lasso Regression	-0.5447
Random Forest	-0.5262

```

Best Model: Random Forest
Test R2 Score: -0.5262
Test RMSE: 9454.4696

print("\n" + "="*50)
print("MODEL JUSTIFICATION")
print("="*50)
print(f"""
The {best_model_name} was selected as the best model for predicting
Turnover based on:

1. **Performance Metrics**:
   - Test R2 Score: {results_df.loc[best_model_name, 'Test R2']:.4f}

```

```

- Test RMSE: {results_df.loc[best_model_name, 'Test RMSE']:.4f}
- Test MAE: {results_df.loc[best_model_name, 'Test MAE']:.4f}

2. Feature Selection Rationale:
- Used correlation analysis to identify most relevant features
- Included lag features to capture temporal dependencies
- Added moving averages to capture trends
- Incorporated volatility measures for risk assessment

3. Model Advantages:
- {'Handles non-linear relationships and feature interactions' if
best_model_name == 'Random Forest' else 'Provides interpretable linear
relationships'}
- {'Robust to outliers and overfitting with ensemble approach' if
best_model_name == 'Random Forest' else 'Simple and computationally
efficient'}
- Good generalization on unseen data (test set)
"""

```

```

=====
MODEL JUSTIFICATION
=====

```

The Random Forest was selected as the best model for predicting Turnover based on:

```

1. Performance Metrics:
- Test R2 Score: -0.5262
- Test RMSE: 9454.4696
- Test MAE: 7571.0970

2. Feature Selection Rationale:
- Used correlation analysis to identify most relevant features
- Included lag features to capture temporal dependencies
- Added moving averages to capture trends
- Incorporated volatility measures for risk assessment

3. Model Advantages:
- Handles non-linear relationships and feature interactions
- Robust to outliers and overfitting with ensemble approach
- Good generalization on unseen data (test set)

```

## 3.1 Augmented Dickey-Fuller (ADF) and KPSS Tests

**Stationarity** is crucial for time series analysis. A stationary time series has:

- Constant mean over time

- Constant variance over time
- Covariance depends only on the lag, not on time

**ADF Test:** Tests for unit root (null hypothesis: series has unit root, i.e., non-stationary) **KPSS**

**Test:** Tests for stationarity (null hypothesis: series is stationary)

```
# Stationarity Testing Functions
def check_stationarity(timeseries, title):
    """
    Perform ADF and KPSS tests for stationarity
    """
    print(f"\n{'='*60}")
    print(f"STATIONARITY TEST RESULTS FOR: {title}")
    print(f'{'='*60}')

    # Augmented Dickey-Fuller test
    print("Augmented Dickey-Fuller Test:")
    adf_result = adfuller(timeseries.dropna())
    print(f'ADF Statistic: {adf_result[0]:.6f}')
    print(f'p-value: {adf_result[1]:.6f}')
    print('Critical Values:')
    for key, value in adf_result[4].items():
        print(f'\t{key}: {value:.3f}')

    if adf_result[1] <= 0.05:
        print("✓ ADF Test: Reject null hypothesis - Time series is STATIONARY")
        adf_stationary = True
    else:
        print("x ADF Test: Fail to reject null hypothesis - Time series is NON-STATIONARY")
        adf_stationary = False

    # KPSS test
    print("\nKwiatkowski-Phillips-Schmidt-Shin Test:")
    kpss_result = kpss(timeseries.dropna(), regression='c')
    print(f'KPSS Statistic: {kpss_result[0]:.6f}')
    print(f'p-value: {kpss_result[1]:.6f}')
    print('Critical Values:')
    for key, value in kpss_result[3].items():
        print(f'\t{key}: {value:.3f}')

    if kpss_result[1] >= 0.05:
        print("✓ KPSS Test: Fail to reject null hypothesis - Time series is STATIONARY")
        kpss_stationary = True
    else:
        print("x KPSS Test: Reject null hypothesis - Time series is NON-STATIONARY")
```

```

        kpss_stationary = False

    # Final interpretation
    print(f"\n{'='*60}")
    if adf_stationary and kpss_stationary:
        print("CONCLUSION: Time series is STATIONARY (both tests agree)")
        return True
    elif not adf_stationary and not kpss_stationary:
        print("CONCLUSION: Time series is NON-STATIONARY (both tests agree)")
        return False
    else:
        print("CONCLUSION: Tests disagree - further investigation needed")
        print("    Consider trend stationarity vs difference stationarity")
        return False

# Test stationarity for the target variable (Turnover)
print("Testing stationarity for the target variable and key features...")

is_stationary = check_stationarity(df_clean[target_col], target_col)

# Test stationarity for other key variables
key_variables = top_features[:3] # Test first 3 original features
stationarity_results = {target_col: is_stationary}

for var in key_variables:
    stationarity_results[var] = check_stationarity(df_clean[var], var)

Testing stationarity for the target variable and key features...

=====
STATIONARITY TEST RESULTS FOR: Turnover (? Cr)
=====
Augmented Dickey-Fuller Test:
ADF Statistic: -3.773713
p-value: 0.003188
Critical Values:
    1%: -3.458
    5%: -2.874
    10%: -2.573
✓ ADF Test: Reject null hypothesis - Time series is STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:
KPSS Statistic: 0.420244
p-value: 0.068429
Critical Values:

```



10%: 0.347  
5%: 0.463  
2.5%: 0.574  
1%: 0.739

✓ KPSS Test: Fail to reject null hypothesis - Time series is STATIONARY

=====  
CONCLUSION: Time series is STATIONARY (both tests agree)

=====  
STATIONARITY TEST RESULTS FOR: Shares Traded  
=====

Augmented Dickey-Fuller Test:

ADF Statistic: -3.714361

p-value: 0.003916

Critical Values:

1%: -3.458  
5%: -2.874  
10%: -2.573

✓ ADF Test: Reject null hypothesis - Time series is STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:

KPSS Statistic: 0.764005

p-value: 0.010000

Critical Values:

10%: 0.347  
5%: 0.463  
2.5%: 0.574  
1%: 0.739

x KPSS Test: Reject null hypothesis - Time series is NON-STATIONARY

=====  
CONCLUSION: Tests disagree - further investigation needed  
Consider trend stationarity vs difference stationarity

=====  
STATIONARITY TEST RESULTS FOR: High  
=====

Augmented Dickey-Fuller Test:

ADF Statistic: -1.893236

p-value: 0.335250

Critical Values:

1%: -3.457  
5%: -2.873  
10%: -2.573

x ADF Test: Fail to reject null hypothesis - Time series is NON-STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:

```
KPSS Statistic: 0.616515
p-value: 0.021135
Critical Values:
    10%: 0.347
    5%: 0.463
    2.5%: 0.574
    1%: 0.739
x KPSS Test: Reject null hypothesis - Time series is NON-STATIONARY
```

```
=====
CONCLUSION: Time series is NON-STATIONARY (both tests agree)
```

```
=====
STATIONARITY TEST RESULTS FOR: Close
=====
```

```
Augmented Dickey-Fuller Test:
ADF Statistic: -1.526882
p-value: 0.520174
Critical Values:
    1%: -3.457
    5%: -2.873
    10%: -2.573
x ADF Test: Fail to reject null hypothesis - Time series is NON-
STATIONARY
```

```
Kwiatkowski-Phillips-Schmidt-Shin Test:
KPSS Statistic: 0.619446
p-value: 0.020869
Critical Values:
    10%: 0.347
    5%: 0.463
    2.5%: 0.574
    1%: 0.739
x KPSS Test: Reject null hypothesis - Time series is NON-STATIONARY
```

```
=====
CONCLUSION: Time series is NON-STATIONARY (both tests agree)
```

## 3.2 Making Time Series Stationary

```
# Apply transformations to make time series stationary
def make_stationary(series, method='differencing'):

    if method == 'differencing':
        return series.diff().dropna()

# Check if any series need to be made stationary
non_stationary_series = [var for var, is_stat in
```

```

stationarity_results.items() if not is_stat]

if non_stationary_series:
    print(f"Non-stationary series found: {non_stationary_series}")

    # Create DataFrame to store transformed series
    df_stationary = df_clean.copy()

    for var in non_stationary_series:
        # Try first differencing
        diff_series = make_stationary(df_clean[var], 'differencing')

        # Test stationarity of differenced series
        print(f"\nTesting stationarity after first differencing for
{var}:")
        is_diff_stationary = check_stationarity(diff_series, f"{var}
(First Difference)")

        if is_diff_stationary:
            df_stationary[f'{var}_diff'] = df_clean[var].diff()
            print(f"✓ First differencing successful for {var}")
        else:
            # Try second differencing
            diff2_series = make_stationary(diff_series,
'differencing')
            print(f"\nTesting stationarity after second differencing
for {var}:")
            is_diff2_stationary = check_stationarity(diff2_series,
f"{var} (Second Difference)")

            if is_diff2_stationary:
                df_stationary[f'{var}_diff2'] =
df_clean[var].diff().diff()
                print(f"✓ Second differencing successful for {var}")
            else:
                print(f"△ Unable to make {var} stationary with
differencing")

```

Non-stationary series found: ['Shares Traded', 'High', 'Close']

Testing stationarity after first differencing for Shares Traded:

```

=====
STATIONARITY TEST RESULTS FOR: Shares Traded (First Difference)
=====
Augmented Dickey-Fuller Test:
ADF Statistic: -6.029236
p-value: 0.000000

```

Critical Values:

1%: -3.459

5%: -2.874

10%: -2.573

✓ ADF Test: Reject null hypothesis - Time series is STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:

KPSS Statistic: 0.373178

p-value: 0.088716

Critical Values:

10%: 0.347

5%: 0.463

2.5%: 0.574

1%: 0.739

✓ KPSS Test: Fail to reject null hypothesis - Time series is STATIONARY

=====

CONCLUSION: Time series is STATIONARY (both tests agree)

✓ First differencing successful for Shares Traded

Testing stationarity after first differencing for High:

=====

STATIONARITY TEST RESULTS FOR: High (First Difference)

=====

Augmented Dickey-Fuller Test:

ADF Statistic: -7.243466

p-value: 0.000000

Critical Values:

1%: -3.458

5%: -2.874

10%: -2.573

✓ ADF Test: Reject null hypothesis - Time series is STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:

KPSS Statistic: 0.139037

p-value: 0.100000

Critical Values:

10%: 0.347

5%: 0.463

2.5%: 0.574

1%: 0.739

✓ KPSS Test: Fail to reject null hypothesis - Time series is STATIONARY

=====

CONCLUSION: Time series is STATIONARY (both tests agree)

✓ First differencing successful for High

Testing stationarity after first differencing for Close:

```
=====
STATIONARITY TEST RESULTS FOR: Close (First Difference)
=====
```

Augmented Dickey-Fuller Test:

ADF Statistic: -7.682519

p-value: 0.000000

Critical Values:

1%: -3.458

5%: -2.874

10%: -2.573

✓ ADF Test: Reject null hypothesis - Time series is STATIONARY

Kwiatkowski-Phillips-Schmidt-Shin Test:

KPSS Statistic: 0.137711

p-value: 0.100000

Critical Values:

10%: 0.347

5%: 0.463

2.5%: 0.574

1%: 0.739

✓ KPSS Test: Fail to reject null hypothesis - Time series is STATIONARY

```
=====
CONCLUSION: Time series is STATIONARY (both tests agree)
```

✓ First differencing successful for Close

## 4.1 Time Series Decomposition and Component Analysis

```
# Time Series Decomposition Analysis
```

```
print("="*60)
```

```
print("TIME SERIES DECOMPOSITION ANALYSIS")
```

```
print("="*60)
```

```
# Perform seasonal decomposition for the target variable
```

```
target_series = df_clean[target_col].dropna()
```

```
# Check if we have enough data points for decomposition
```

```
if len(target_series) >= 24: # Need at least 2 complete cycles
```

```
    # Determine period for decomposition
```

```
    period = min(12, len(target_series) // 4) # Quarterly or monthly seasonality
```

```
    # Perform additive decomposition
```

```
    decomposition_add = seasonal_decompose(target_series,
model='additive', period=period)
```

```

    # Perform multiplicative decomposition
    decomposition_mult = seasonal_decompose(target_series,
model='multiplicative', period=period)

    # Plot decomposition
    fig, axes = plt.subplots(4, 2, figsize=(20, 16))
    fig.suptitle(f'Time Series Decomposition - {target_col}',
    fontsize=16, y=0.98)

    # Additive decomposition
    decomposition_add.observed.plot(ax=axes[0, 0], title='Original
(Additive)')
    decomposition_add.trend.plot(ax=axes[1, 0], title='Trend
(Additive)')
    decomposition_add.seasonal.plot(ax=axes[2, 0], title='Seasonal
(Additive)')
    decomposition_add.resid.plot(ax=axes[3, 0], title='Residual
(Additive)')

    # Multiplicative decomposition
    decomposition_mult.observed.plot(ax=axes[0, 1], title='Original
(Multiplicative)')
    decomposition_mult.trend.plot(ax=axes[1, 1], title='Trend
(Multiplicative)')
    decomposition_mult.seasonal.plot(ax=axes[2, 1], title='Seasonal
(Multiplicative)')
    decomposition_mult.resid.plot(ax=axes[3, 1], title='Residual
(Multiplicative)')

    for ax in axes.flat:
        ax.grid(True, alpha=0.3)
        ax.tick_params(axis='x', rotation=45)

    plt.tight_layout()
    plt.show()

    # Analyze components
    print(f"\nDECOMPOSITION ANALYSIS:")
    print(f"Period used for decomposition: {period}")

    # Trend analysis
    trend_change = decomposition_add.trend.dropna()
    if len(trend_change) > 1:
        overall_trend = "Increasing" if trend_change.iloc[-1] >
trend_change.iloc[0] else "Decreasing"
        print(f"Overall Trend: {overall_trend}")
        print(f"Trend Range: {trend_change.min():.2f} to
{trend_change.max():.2f}")

```

```

    # Seasonality analysis
    seasonal_strength = np.std(decomposition_add.seasonal.dropna()) /
np.std(target_series)
    print(f"Seasonal Strength: {seasonal_strength:.4f}")
    if seasonal_strength > 0.1:
        print("✓ Significant seasonality detected")
        has_seasonality = True
    else:
        print("✗ No significant seasonality detected")
        has_seasonality = False

    # Residual analysis
    residuals = decomposition_add.resid.dropna()
    residual_mean = residuals.mean()
    residual_std = residuals.std()
    print(f"Residual Mean: {residual_mean:.6f}")
    print(f"Residual Std: {residual_std:.6f}")

else:
    print("Insufficient data for seasonal decomposition (need at least
24 observations)")
    has_seasonality = False


    # Simple trend analysis
    if len(target_series) > 1:
        trend_slope = np.polyfit(range(len(target_series)),
target_series, 1)[0]
        print(f"Linear trend slope: {trend_slope:.4f}")
        if abs(trend_slope) > target_series.std() * 0.01:
            print("✓ Trend detected")
        else:
            print("✗ No significant trend detected")

```

```

=====
TIME SERIES DECOMPOSITION ANALYSIS
=====

```



#### DECOMPOSITION ANALYSIS:

Period used for decomposition: 12

Overall Trend: Decreasing

Trend Range: 22776.27 to 40743.09

Seasonal Strength: 0.1240

✓ Significant seasonality detected

Residual Mean: -82.711606

Residual Std: 7405.025007

## 4.2 ARIMA and SARIMA Model Development

### ARIMA vs SARIMA Models:

#### ARIMA (AutoRegressive Integrated Moving Average):

- **Components:** AR(p), I(d), MA(q)
- **Best for:** Non-seasonal time series data
- **Parameters:**



- p: Number of lag observations (autoregressive terms)
- d: Degree of differencing (integrated terms)
- q: Size of moving average window (moving average terms)

### SARIMA (Seasonal ARIMA):

- **Components:** ARIMA + Seasonal components (P, D, Q, s)
- **Best for:** Time series with seasonal patterns
- **Parameters:** ARIMA(p,d,q) × (P,D,Q,s)
  - P, D, Q: Seasonal AR, differencing, and MA terms
  - s: Length of seasonal cycle

### When to prefer each:

- **ARIMA:** When data shows no clear seasonal patterns
- **SARIMA:** When data exhibits seasonal behavior (daily, weekly, monthly patterns)

```
# ARIMA and SARIMA Model Training and Evaluation
print("="*60)
print("ARIMA AND SARIMA MODEL TRAINING")
print("="*60)

# Split data for time series forecasting
train_size = int(len(target_series) * 0.8)
train_data = target_series[:train_size]
test_data = target_series[train_size:]

print(f"Training data size: {len(train_data)}")
print(f"Test data size: {len(test_data)}")

# Initialize results dictionary
ts_results = {}

# 1. ARIMA Model
print("\n1. Building ARIMA Model...")
try:
    arima_model = ARIMA(train_data, order=arima_order)
    arima_fitted = arima_model.fit()

    # Make predictions
    arima_forecast = arima_fitted.forecast(steps=len(test_data))
    arima_forecast_ci =
    arima_fitted.get_forecast(steps=len(test_data)).conf_int()

    # Calculate metrics
    arima_mae = mean_absolute_error(test_data, arima_forecast)
    arima_rmse = np.sqrt(mean_squared_error(test_data,
    arima_forecast))
```

```

ts_results['ARIMA'] = {
    'model': arima_fitted,
    'forecast': arima_forecast,
    'conf_int': arima_forecast_ci,
    'mae': arima_mae,
    'rmse': arima_rmse,
    'aic': arima_fitted.aic,
    'params': arima_order
}

print(f"✓ ARIMA{arima_order} Model Summary:")
print(f"  AIC: {arima_fitted.aic:.4f}")
print(f"  MAE: {arima_mae:.4f}")
print(f"  RMSE: {arima_rmse:.4f}")

except Exception as e:
    print(f"✗ ARIMA model failed: {e}")

# 2. SARIMA Model (if seasonality detected)
if has_seasonality and len(target_series) >= 24:
    print("\n2. Building SARIMA Model...")

    # Determine seasonal parameters
    seasonal_period = min(12, len(target_series) // 4)
    sarima_seasonal_order = (1, 1, 1, seasonal_period) # Simple
    seasonal parameters

    try:
        sarima_model = SARIMAX(train_data, order=arima_order,
                                seasonal_order=sarima_seasonal_order)
        sarima_fitted = sarima_model.fit(dispatch=False)

        # Make predictions
        sarima_forecast = sarima_fitted.forecast(steps=len(test_data))
        sarima_forecast_ci =
        sarima_fitted.get_forecast(steps=len(test_data)).conf_int()

        # Calculate metrics
        sarima_mae = mean_absolute_error(test_data, sarima_forecast)
        sarima_rmse = np.sqrt(mean_squared_error(test_data,
        sarima_forecast))

        ts_results['SARIMA'] = {
            'model': sarima_fitted,
            'forecast': sarima_forecast,
            'conf_int': sarima_forecast_ci,
            'mae': sarima_mae,
            'rmse': sarima_rmse,
            'aic': sarima_fitted.aic,

```

```

        'params': (arima_order, sarima_seasonal_order)
    }

    print(f"✓ SARIMA{arima_order}x{sarima_seasonal_order} Model
Summary:")
    print(f"    AIC: {sarima_fitted.aic:.4f}")
    print(f"    MAE: {sarima_mae:.4f}")
    print(f"    RMSE: {sarima_rmse:.4f}")

    except Exception as e:
        print(f"x SARIMA model failed: {e}")
else:
    print("\n2. SARIMA Model skipped (no significant seasonality or
insufficient data)")

# 3. Simple baseline models for comparison
print("\n3. Building Baseline Models...")

# Naive forecast (last value)
naive_forecast = [train_data.iloc[-1]] * len(test_data)
naive_mae = mean_absolute_error(test_data, naive_forecast)
naive_rmse = np.sqrt(mean_squared_error(test_data, naive_forecast))

ts_results['Naive'] = {
    'forecast': naive_forecast,
    'mae': naive_mae,
    'rmse': naive_rmse
}

# Moving average forecast
ma_window = min(5, len(train_data) // 4)
ma_forecast = [train_data.rolling(window=ma_window).mean().iloc[-1]] *
len(test_data)
ma_mae = mean_absolute_error(test_data, ma_forecast)
ma_rmse = np.sqrt(mean_squared_error(test_data, ma_forecast))

ts_results['Moving Average'] = {
    'forecast': ma_forecast,
    'mae': ma_mae,
    'rmse': ma_rmse
}

print(f"✓ Naive Forecast - MAE: {naive_mae:.4f}, RMSE:
{naive_rmse:.4f}")
print(f"✓ Moving Average Forecast - MAE: {ma_mae:.4f}, RMSE:
{ma_rmse:.4f}")

# Compare models
print("\n" + "="*60)
print("MODEL COMPARISON")

```

```

print("="*60)

comparison_df = pd.DataFrame({
    'Model': list(ts_results.keys()),
    'MAE': [ts_results[model]['mae'] for model in ts_results.keys()],
    'RMSE': [ts_results[model]['rmse'] for model in ts_results.keys()]
})

print(comparison_df.round(4))

# Find best model
best_ts_model = comparison_df.loc[comparison_df['MAE'].idxmin(),
    'Model']
print(f"\nBest Time Series Model: {best_ts_model}")
print(f"Best MAE: {ts_results[best_ts_model]['mae']:.4f}")
print(f"Best RMSE: {ts_results[best_ts_model]['rmse']:.4f}")

```

## ARIMA AND SARIMA MODEL TRAINING

Training data size: 199

Test data size: 50

1. Building ARIMA Model...

x ARIMA model failed: name 'arima\_order' is not defined

2. Building SARIMA Model...

x SARIMA model failed: name 'arima\_order' is not defined

3. Building Baseline Models...

✓ Naive Forecast - MAE: 6941.9970, RMSE: 8457.9343

✓ Moving Average Forecast - MAE: 8376.3540, RMSE: 9751.8257

## MODEL COMPARISON

	Model	MAE	RMSE
0	Naive	6941.997	8457.9343
1	Moving Average	8376.354	9751.8257

Best Time Series Model: Naive

Best MAE: 6941.9970

Best RMSE: 8457.9343